

```
# Receiver Operating Characteristic (ROC) curves
# examines the trade-off between the detection of true positives,
# while avoiding the false positives.
# graph between true pos rate (sensitivity) & false pos rate (1-specificity)
```

```
dt2<-read.csv(file.choose(), stringsAsFactors = F)
dt<-dt2
dt[c(2,3,6,7,9,13,14)]<-data.frame(lapply(dt[c(2,3,6,7,9,13,14)],factor))
dt <- dt[c(2,3,9,10,11,12,14)]
```

```
dummysize <- function(i, exclude=c("")){
  x<-dt[i]
  if(is.factor(dt[[i]]) & !identical(names(x) %in% exclude, TRUE)){
    x<-data.frame(model.matrix(~.-1,data=x))[-1]
  }
  return(x)
}
normalize<-function(x){
  if(is.numeric(x)) x <- (x-min(x))/(max(x)-min(x))
  return (x)
}
```

```
dt <- data.frame(lapply(seq_along(dt),dummysize, "statlogheart"))
str(dt)
```

```
#Many models assume that the 2nd label is positive event
#By default, Have disease can be 1st label because of chronological order
# hence we will set it as 2nd label
dt$statlogheart <- factor(dt$statlogheart, levels=c("1", "2"),
  labels=c("No disease", "Have disease"))
```

```
dt<-data.frame(lapply(dt, normalize))
str(dt)
```

```
rows <- nrow(dt)
train.index <- sample(1:rows, floor(0.7*rows), replace=FALSE)
dt.train <- dt[train.index,]
dt.test <- dt[-train.index,]
```

```
#####knn
library(class)
# install.packages('ROCR')
library(ROCR)
library(gmodels)
```

```
#if you have predicted probabilities for the class label, say pred.prob, and you have class
label, how would you calculated predicted probability for positive event.
```

```
# eg
# pred.class pred.prob pred.pos.prob
# pos      0.6      0.6
# neg      0.7      1-0.7 = 0.3
# neg      0.8      1-0.8 = 0.2
# pos      0.65     0.65
# pred.pos.prob <- ifelse(pred.class=="pos", pred.prob, 1-pred.prob)
```

```
class.knn.test <- knn(dt.train[-9],dt.test[-9],dt.train$statlogheart,k=3, prob=TRUE)
prob.knn.test <- ifelse(class.knn.test=="Have disease", attr(class.knn.test,"prob"), 1-
attr(class.knn.test,"prob"))
```

```
#to check if pred.knn.prob is correct we recalculate class labels and compare them with line
pred.knn.class
head (prob.knn.test)
# idx <- which(pred.knn.prob==0.5)
# pred.knn.prob[15]
class.check <- ifelse(prob.knn.test>0.5, "Have disease", "No disease")
table(class.knn.test,class.check)
#since there were no errors, prob calculated above for positive event is correct
CrossTable(dt.test$statlogheart, class.knn.test)
```

```

# Every evaluation method such as lift, and ROC in ROCR needs a prediction object.
# It transforms the input data into a standardized format needed by ROCR package.
# 1st parameter for prediction is prob of positive event
# 2nd parameter is the actual class label
# 3rd parameter is label ordering
# NOTE: any time you are going to plot ROC curves, don't forget to set the
# order of class labels of your target variable. Set 1st level as negative event
# and 2nd level as positive event
predObj.knn.test <- prediction(prob.knn.test, dt.test$statlogheart,
                              label.ordering = c("No disease", "Have disease"))

#lift
lift.knn.test <- performance(predObj.knn.test, measure = "lift", x.measure = "rpp")
plot(lift.knn.test, main="Lift curve for KNN model")

#Cumulative gain is more informative than lift
#cumulative gain
gain.knn.test <- performance(predObj.knn.test, "tpr", "rpp")
plot(gain.knn.test, main="Cumulative gain curve for KNN model")

#roc
roc.knn.test <- performance(predObj.knn.test, "tpr", "fpr")
# add a reference line to the graph
# set the intercept to a = 0 and the slope to b = 1,
# lwd adjusts the line thickness, while the lty adjusts the type of line
# worst classifier
plot(roc.knn.test, main="ROC curve for KNN model", col = "khaki", lwd = 2)
# worst classifier
abline(a = 0, b = 1, lwd = 2, lty = 2)

#Area under the curve
auc.knn.test <- performance(predObj.knn.test, "auc")
#To get AUC, unlist auc.knn.test object, which is a list
unlist(auc.knn.test@y.values)

#####NB
library(e1071)
#NOTE: other models can take target as num, but bays has to take it as factor
model.nb <- naiveBayes(statlogheart~, data=dt.train, laplace=1)
class.nb.test <- predict(model.nb, dt.test)
CrossTable(dt.test$statlogheart, class.nb.test)

prob.nb.test <- predict(model.nb, dt.test, type="raw")
# output of type raw gives probabilities for both levels
head(prob.nb.test)
# since the prob for positive event is in 2nd column, & further analysis uses
# only positive event probabilities, we need to subset prob.nb.test
prob.nb.test <- prob.nb.test[,2]

class.check <- ifelse(prob.nb.test>0.5, "Have disease", "No disease")
CrossTable(class.nb.test, class.check)
#predicted.nb

#calculate prediction object for various methods in ROCR package
predObj.nb.test <- prediction(prob.nb.test, dt.test$statlogheart,
                              label.ordering = c("No disease", "Have disease"))

#Lift
lift.nb.test <- performance(predObj.nb.test, "lift", "rpp")
plot(lift.nb.test, main="Lift curve for Naive Bayes model")

#Cumulative gain
gain.nb.test <- performance(predObj.nb.test, "tpr", "rpp")
plot(gain.nb.test, main="Cumulative gain curve for Naive Bayes model")

#ROC
roc.nb.test <- performance(predObj.nb.test, "tpr", "fpr")
plot(roc.nb.test, main="ROC curve for Naive Bayes model", col = "navy blue", lwd = 2)
# worst classifier

```

```
abline(a = 0, b = 1, lwd = 2, lty = 2)
```

```
#Area under the curve
```

```
auc.nb.test <- performance(predObj.nb.test, "auc")
```

```
#To get AUC, unlist auc.knn.test object, which is a list
```

```
unlist(auc.nb.test@y.values)
```

```
#####logit
```

```
model.logit <- glm(statlogheart~.,data=dt.train, family = binomial)
```

```
summary(model.logit)
```

```
#type="link" odds ratio
```

```
#type="response" probabilities
```

```
prob.logit.test <- predict(model.logit,dt.test, type="response")
```

```
# output is probability for positive event
```

```
class.logit.test <- ifelse(prob.logit.test>0.5,"Have disease","No disease")
```

```
CrossTable(dt.test$statlogheart,class.logit.test)
```

```
predObj.logit.test <- prediction(prob.logit.test,dt.test$statlogheart,  
                                label.ordering = c("No disease", "Have disease"))
```

```
#lift
```

```
lift.logit.test <- performance(predObj.logit.test, "lift", "rpp")
```

```
plot(lift.logit.test, main="Lift curve for logit model")
```

```
#cumulative gain
```

```
gain.logit.test <- performance(predObj.logit.test, "tpr", "rpp")
```

```
plot(gain.logit.test, main="Cumulative gain curve for logit model")
```

```
#roc
```

```
roc.logit.test<-performance(predObj.logit.test, "tpr", "fpr")
```

```
plot(roc.logit.test, main="ROC curve for logit model", col = "lavender", lwd = 2)
```

```
# worst classifier
```

```
abline(a = 0, b = 1, lwd = 2, lty = 2)
```

```
#Area under the curve
```

```
auc.logit.test <- performance(predObj.logit.test, "auc")
```

```
#To get AUC, unlist auc.logit.test object, which is a list
```

```
unlist(auc.logit.test@y.values)
```

```
#####decision tree
```

```
library(C50)
```

```
model.tree <- C5.0(statlogheart ~.,data=dt.train)
```

```
summary(model.tree)
```

```
class.tree.test <- predict(model.tree, dt.test)
```

```
prob.tree.test <- predict(model.tree,dt.test,type="prob")
```

```
head(prob.tree.test)
```

```
prob.tree.test <- prob.tree.test[,2]
```

```
CrossTable(dt.test$statlogheart,class.tree.test)
```

```
predObj.tree.test <- prediction(prob.tree.test,dt.test$statlogheart,  
                                label.ordering = c("No disease", "Have disease"))
```

```
#lift
```

```
lift.tree.test <- performance(predObj.tree.test, "lift", "rpp")
```

```
plot(lift.tree.test, main="Lift curve for tree model")
```

```
#cumulative gain
```

```
gain.tree.test <- performance(predObj.tree.test, "tpr", "rpp")
```

```
plot(gain.tree.test, main="Cumulative gain curve for tree model")
```

```
#roc
```

```
roc.tree.test<-performance(predObj.tree.test, "tpr", "fpr")
```

```
plot(roc.tree.test, main="ROC curve for tree model", col = "turquoise", lwd = 2)
```

```
# worst classifier
```

```
abline(a = 0, b = 1, lwd = 2, lty = 2)
```

```
#Area under the curve
```

```
auc.tree.test <- performance(predObj.tree.test, "auc")
```

```
#To get AUC, unlist auc.tree.test object, which is a list
```

```
unlist(auc.tree.test@y.values)
```

SVM

```
library(kernlab)
model.svm <- ksvm(statlogheart ~., data=dt.train, kernel="rbfdot", prob.model=T)
class.svm.test <- predict(model.svm, dt.test)
head(class.svm.test)
prob.svm.test <- predict(model.svm, dt.test, type="probabilities")
head(prob.svm.test)
prob.svm.test <- prob.svm.test[,2]
class.check <- ifelse(prob.svm.test>0.5, "Have disease", "No disease")
table(class.svm.test,class.check)
CrossTable(dt.test$statlogheart,class.svm.test)
class(dt.test$statlogheart)
predObj.svm.test <- ROCR::prediction(prob.svm.test,dt.test$statlogheart,
label.ordering = c("No disease", "Have disease"))
```

#lift

```
lift.svm.test <- performance(predObj.svm.test, "lift", "rpp")
plot(lift.svm.test, main="Lift curve for svm model")
```

#cumulative gain

```
gain.svm.test <- performance(predObj.svm.test, "tpr", "rpp")
plot(gain.svm.test, main="Cumulative gain curve for svm model")
```

#roc

```
roc.svm.test<-performance(predObj.svm.test, "tpr", "fpr")
plot(roc.svm.test, main="ROC curve for svm model", col = "salmon", lwd = 2)
# worst classifier
abline(a = 0, b = 1, lwd = 2, lty = 2)
```

#Area under the curve

```
auc.svm.test <- performance(predObj.svm.test, "auc")
#To get AUC, unlist auc.svm.test object, which is a list
unlist(auc.svm.test@y.values)
```

Neuralnet

```
library(neuralnet)
model.nn <- neuralnet(statlogheart ~., data=dt.train, hidden = 1)
computed.nn <- compute(model.nn, dt.test)
#NOTE: neuralnet package sorts labels chronologically
head(computed.nn$net.result)
prob.nn.test <- computed.nn$net.result
head(prob.nn.test)
prob.nn.test <- (prob.nn.test[,1])
head(prob.nn.test)

class.nn.test <- ifelse(prob.nn.test>0.5, "Have disease", "No disease")
CrossTable(dt.test$statlogheart,class.nn.test)
summary(prob.nn.test)
predObj.nn.test <- ROCR::prediction(prob.nn.test,dt.test$statlogheart,
label.ordering = c("No disease", "Have disease"))
```

#lift

```
lift.nn.test <- performance(predObj.nn.test, "lift", "rpp")
plot(lift.nn.test, main="Lift curve for nn model")
```

#cumulative gain

```
gain.nn.test <- performance(predObj.nn.test, "tpr", "rpp")
plot(gain.nn.test, main="Cumulative gain curve for nn model")
```

#roc

```
roc.nn.test<-performance(predObj.nn.test, "tpr", "fpr")
plot(roc.nn.test, main="ROC curve for nn model", col = "orange", lwd = 2)
# worst classifier
abline(a = 0, b = 1, lwd = 2, lty = 2)
```

#Area under the curve

```
auc.nn.test <- performance(predObj.nn.test, "auc")
#To get AUC, unlist auc.nn.test object, which is a list
unlist(auc.nn.test@y.values)
```

```
##ROC curves on same plot
plot(roc.nb.test, xlab="False Positive Rate", ylab="True Positive Rate",
     type='l', col = "navy blue")
plot(roc.tree.test, col = "turquoise", add=T)
plot(roc.knn.test, col = "khaki", add=T)
plot(roc.logit.test, col = "lavender", add=T)
plot(roc.svm.test, col = "salmon", add=T)
plot(roc.nn.test, col = "green", add=T)
abline(v=0,lty=2, lwd=1, col="gray60")
abline(h=1,lty=2, lwd=1, col="gray60")
abline(a=0,b=1,lty=2, lwd=1, col="gray60")

##ROC curves on same plot using ggplot2()
library(ggplot2)
ggplot()+
  geom_line(data=data.frame(x=roc.tree.test@x.values[[1]], y=roc.tree.test@y.values[[1]]),
            aes(x=x, y=y, col="red")) +
  geom_line(data=data.frame(x=roc.svm.test@x.values[[1]], y=roc.svm.test@y.values[[1]]),
            aes(x=x, y=y, col="blue"))
```

```
##ROC curves on same plot
plot(xlab="False Positive Rate", ylab="True Positive Rate",
     roc.nb.test @x.values[[1]], roc.nb.test @y.values[[1]],type='l', col = "navy blue")
lines(roc.tree.test@x.values[[1]], roc.tree.test@y.values[[1]], col = "turquoise")
lines(roc.knn.test@x.values[[1]], roc.knn.test@y.values[[1]], col = "khaki")
lines(roc.logit.test@x.values[[1]], roc.logit.test@y.values[[1]], col = "lavender")
lines(roc.svm.test@x.values[[1]], roc.svm.test@y.values[[1]], col = "salmon")
lines(roc.nn.test@x.values[[1]], roc.nn.test@y.values[[1]], col = "green")
abline(v=0,lty=2, lwd=1, col="gray60")
abline(h=1,lty=2, lwd=1, col="gray60")
abline(a=0,b=1,lty=2, lwd=1, col="gray60")
```

```
lift.knn.test@y.values[[1]] <- ifelse(lift.knn.test@y.values[[1]]=="NaN",
0, lift.knn.test@y.values[[1]])
```

#####Chapter 10

Confusion matrixes in R ----

use file sms_results.csv

dt<- read.csv(file.choose(),stringsAsFactors = T)

the first several test cases

head(dt)

test cases where the model is less confident say between 0.4 and 0.6

dt[dt\$prob_spam > 0.4 & dt\$prob_spam < 0.6,]

head(subset(dt, prob_spam > 0.40 & prob_spam < 0.60))

test cases where the model was wrong

head(subset(dt, actual_type != predict_type))

tabulate results

table(dt\$actual_type, dt\$predict_type)

#using CrossTable

library(gmodels)

CrossTable(dt\$actual_type, dt\$predict_type)

alternative solution using the formula interface (not shown in book)

xtabs(~ actual_type + predict_type, sms_results)

using the CrossTable function

accuracy and error rate calculation --

accuracy = TP+TN/Total

(152 + 1203) / (152 + 1203 + 4 + 31)

error rate = FP+FN/total

(4 + 31) / (152 + 1203 + 4 + 31)

error rate = 1 - accuracy

1 - 0.9748201

Beyond accuracy: other performance measures ----

install.packages('caret')

library(caret)

confusionMatrix(sms_results\$predict_type, sms_results\$actual_type, positive = "spam")

CrossTable(sms_results\$predict_type, sms_results\$actual_type)

Kappa statistic adjusts accuracy by accounting for the possibility that

correct prediction can happen by chance alone

pr of accuracy = TP/total + TN/total

pr_a <- 0.865 + 0.109

pr_a

#pr of being correct by chance alone = PP*AP/total^2 + PN*AN/total^2

pr_e <- 0.868 * 0.888 + 0.132 * 0.112

pr_e

#k= Pr(a)-Pr(e)/1-Pr(e)

k <- (pr_a - pr_e) / (1 - pr_e)

k

calculate kappa via the vcd package

install.packages('vcd')

library(vcd)

Kappa(table(sms_results\$actual_type, sms_results\$predict_type))

calculate kappa via the irr package

install.packages('irr')

library(irr)

kappa2(sms_results[1:2])\$value

Sensitivity

Sensitivity proportion of positive examples that were correctly classified

Sensitivity (true positive rate) = TP/TP+FN

sens <- 152 / (152 + 31)

sens

```

# Specificity
# measures the proportion of negative examples that were correctly classified.
# Specificity (true negative rate) = TN/TN+FP
spec <- 1203 / (1203 + 4)
spec

# example using the caret package
library(caret)
sensitivity(sms_results$predict_type, sms_results$actual_type, positive = "spam")
specificity(sms_results$predict_type, sms_results$actual_type, negative = "ham")

# Precision (positive predictive value) = TP/TP+FP
# is the proportion of positive examples that are truly positive;
# Google search results
prec <- 152 / (152 + 4)
prec

# Recall = Sensitivity = TP/TP+FN
# is a measure of how complete the results are.
rec <- 152 / (152 + 31)
rec

# example using the caret package
library(caret)
posPredValue(sms_results$predict_type, sms_results$actual_type, positive = "spam")
sensitivity(sms_results$predict_type, sms_results$actual_type, positive = "spam")

# F-measure = 2*Pre*Rec/Pre+Rec Harmonic mean of Pre & Rec
# harmonic mean is used rather than the common arithmetic mean
# since both precision and recall are expressed as proportions between zero and one,
# which can be interpreted as rates.
# eg average rate of a bucket getting filled when there are 2 taps
f <- (2 * prec * rec) / (prec + rec)
f

f <- (2 * 152) / (2 * 152 + 4 + 31)
f

## Visualizing Performance Tradeoffs ----
install.packages('ROCR')
library(ROCR)
pred <- prediction(predictions = sms_results$prob_spam,
labels = sms_results$actual_type)

# Receiver Operating Characteristic (ROC) curves
# examines the trade-off between the detection of true positives,
# while avoiding the false positives.
# graph between true pos rate (sensitivity) & false pos rate (1-specificity)
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
plot(perf, main = "ROC curve for SMS spam filter", col = "blue", lwd = 2)

# add a reference line to the graph
# set the intercept to a = 0 and the slope to b = 1,
# lwd adjusts the line thickness, while the lty adjusts the type of line
# worst classifier
abline(a = 0, b = 1, lwd = 2, lty = 2)
# perfect classifier
abline(a = 1, b = 0, lwd = 2, lty = 2, col='red')
abline(v=0, lwd=2, lty=2, col='red')

# calculate AUC
# The closer the curve is to the perfect classifier,
# the better it is at identifying positive values
perf.auc <- performance(pred, measure = "auc")
# perf.auc is an R object (specifically known as an S4 object),
# use a special type of notation to access the values stored within.
# S4 objects hold information in positions known as slots.
# str() function can be used to see all of an object's slots:
str(perf.auc)
# to see value in specific slot we can use @
# further to get value and not as list you can either use unlist or [[]]
auc <- perf.auc@y.values[[1]]

```

```
auc
unlist(perf.auc@y.values)
```

```
## Estimating Future Performance ----
```

```
# partitioning data
library(caret)
credit <- read.csv("credit.csv")
```

```
# Holdout method-- so far into training and testing
# using test for selecting model does not make it
# an unbiased measure of the performance on unseen data
rows <- nrow(credit)
set.seed(7)
train.size <- floor(rows*0.6)
val.size <- floor(rows*0.2)
```

```
train.vec <- sample(rows, train.size)
# get validation vector: from all rows index remove train.vec
val.vec <- sample (data.frame(1:rows)[-train.vec,], val.size)
train.set <- credit[train.vec,]
val.set <- credit[val.vec,]
test.set <- credit[-c(train.vec,val.vec),]
```

```
#other way of splitting
random_ids <- order(runif(rows))
credit_train <- credit[random_ids[1:train.size],]
credit_validate <- credit[random_ids[(train.size+1):(train.size+val.size)], ]
credit_test <- credit[random_ids[(train.size+val.size+1):1000], ]
```

```
## using caret function
# in_train <- createDataPartition(credit$default, p = 0.75, list = FALSE)
# credit_train <- credit[in_train, ]
# credit_test <- credit[-in_train, ]
```

```
# Stratified sampling: each partition may have a larger or smaller proportion of some classes
```

```
# Repeated holdout method: k-fold CV
# 10-fold CV
```

```
folds <- createFolds(credit$default, k = 10)
str(folds)
credit01_train <- credit[-folds$Fold01, ]
credit01_test <- credit[folds$Fold01, ]
```

```
## Automating 10-fold CV for a C5.0 Decision Tree using lapply() ----
library(caret)
library(C50)
library(irr)
```

```
credit <- read.csv("credit.csv")
```

```
set.seed(123)
folds <- createFolds(credit$default, k = 10)
```

```
#let us try to run 10-fold CV on decision tree
library(C50)
#if you have to do this manually then for the first iteration following is what you have to do
train.set<-dt[-folds$Fold01,]
test.set <- dt[folds$Fold01,]
```

```
foldAnalysis <- function(x){
  train.set <- dt[-x,]
  test.set <- dt[x,]
  model<- C5.0(default~., data=train.set)
  predicted<-predict(model,test.set)
  kappa<-kappa2(data.frame(test.set$default,predicted))$value
  return (kappa)
}
allFolds<-unlist(lapply(folds,foldAnalysis))
allFolds
#following gives you mean kappa using 10-fold cross validation
```



```
mean(allFolds)
```

CHAPTER 11

```
#Extract credit.csv from the class data
```

```
dt <- read.csv(file.choose())
```

```
str(dt)
```

```
## Classification and Regression Training (CART) ##
```

```
install.packages("caret")
```

```
library(caret)
```

```
modelLookup("knn")
```

```
modelLookup("C5.0")
```

```
model <- train(default ~. , data=dt, method="C5.0")
```

```
model
```

```
#instead of it using bootstrap sampling by default, we want caret to use
```

```
# 10-fold cross validation and the selection function should not be the
```

```
# model with best performance over a metric, rather a model that is within
```

```
# one standard error of the best performing model and less complex
```

```
trCtrl <- trainControl(method="cv", n=10, selectionFunction = "oneSE")
```

```
# instead of relying on caret to pick random 3 values of parameters, we
```

```
# specified values it should try for running models.
```

```
pRange <- expand.grid(.model="tree", .winnow=F, .trials=c(1,5,10,15,30,50))
```

```
model <- train(default ~., data=dt, method="C5.0",
```

```
metric= "Kappa",
```

```
trCtrl=trCtrl,
```

```
tuneGrid=pRange)
```

```
model
```

```
predicted <- predict(model, dt.test)
```

```
## Ensembles combine various weak learners to create a strong learner
```

```
# Weak learner is a model that has error rate slightly less than 0.5
```

```
## Bootstrap aggregation (Bagging)
```

```
# You give equal weight to the prediction from various models
```

```
## Adaboost (Adaptive boosting)
```

```
# Steps:
```

```
# 1) In the first iteration all records get equal weighted. Based on training
```

```
# set you will estimate a model, and make predictions for the entire data,
```

```
# and calculate errors;
```

```
# 2) Errors that are bigger will get higher weight in the next iteration;
```

```
# You will run say 20 models, and calculate misclassification for each of
```

```
# the 20 models. But the weight while calculating misclassification the
```

```
# weight will not be the same across instances. So you will pick a model
```

```
# that gives more importance to records that have higher error in them.
```

REGULARIZATION

```
In regression models, in order to reduce variance we introduce some bias in the model.
```

```
Instead of minimizing just Sum of square of errors, we also add a penalty on coefficients.
```

```
#Ridge regression
```

```
# Specifically, in case of Ridge regression, we add L2 norm of the coefficients of features.
```

```
# L2 norm is sum of square of elements of a vector
```

```
# Objective: Min SSE + lambda * sum of square of coefficients
```

```
# This is equivalent to adding the constraint of a circle in the case of 2 features;
```

```
# Hence, the coefficients will never be zero.
```

```
# They may be very small.
```

```
# If lambda is high then you reduce the risk of overfitting;
```

```
#Lasso regression
```

```
# we add L1 norm of the coefficients of features;
```

```
# L1 norm is sum of absolute value of elements of a vector;
```

```
# Objective: Min SSE + lambda * sum of absolute value of coefficients
```

```
# This is equivalent to adding the constraint of a rhombus in the case of 2 features.
```

```
# Hence, the coefficients can be zero, and this can help in feature selection.
```

Elastic net regression
we add linear combination of L2 and L1 norm of the coefficients of features;
objective: Min SSE + lambda * ((1-alpha)*sum of square of coefficients + alpha* abs value of coefficients)
if alpha =1 => Lasso regression
if alpha =0 => Ridge regression
0 <= alpha <= 1

XGBoost is one of the most prominent machine learning technique and has been a top choice by people winning
Kaggle competitions. It is pretty versatile and can be used for both classification, numeric prediction and ranking
It improves GBM by 10X.

*****After this, you may not worry about it for the exam*****
Good balanced reading: <https://www.datacamp.com/community/tutorials/tutorial-ridge-lasso-elastic-net>
Good reading for Ridge, Lasso and Elasticnet R code is https://web.stanford.edu/~hastie/glmnet/glmnet_alpha.html
Good reading for maths behind Ridge, Lasso and Elasticnet is <http://statweb.stanford.edu/~tibs/sta305files/Rudyregularization.pdf>
For graphical explanation see this:
[https://en.wikipedia.org/wiki/Regularization_\(mathematics\)#Regularizers_for_sparsity](https://en.wikipedia.org/wiki/Regularization_(mathematics)#Regularizers_for_sparsity)

Gradient Boosting Machine (GBM) is like adaptive boosting
it can be used for classification, numeric prediction / regression, ranking data
GBM: Gradient descent is partial of Loss function wrt to predicted target;
Friedman (1999) showed that in case of MSE (L2 loss), gradient is equal to the mean of residuals,
while in MAE (L1 loss) the gradient is the sign of residuals;
$F_m(x) = F_{m-1}(x) + \rho * h_m(x)$; where rho is the learning rate;
where $h_m(x) = y - F_{m-1}(x)$

Steps:
1) Take the median (mean) value of target for the MAE (MSE) loss. $F_0(x) = \text{median (mean) of } X$; Calculate residuals;
2) Take the sign (values) of residuals as target, and find a split that minimizes the variance in
sign (actual) values of residuals. Residual values is called direction vector, which means both sign & value;
3) Take the prediction in each leaf as the median (mean) residual. The sign vector is used for grouping/splitting
purposes in L1 loss, but the actual prediction is in fact a residual, just like it is for the L2 loss algorithm.
4) Multiply this predicted error with learning rate, and add it into summed predicted value;
5) Find the new residual by subtracting summed predicted value from the actual target;
Repeat step 2

XGBoost is an ensemble learning method;
It implements parallel processing, and is hence fast (typically 10X faster than GBM)
it implements regularisation helping reduce overfit
it generally has more predictive power than traditional Gradient boosting method (GBM) or Adaboost.
While Gradient Boosting uses greedy approach to try various models, XGBoost uses Taylor expansion till 2nd order
to calculate the optimum weight for tree leaves.
XGBoost doesn't explore all possible tree structures but builds a tree greedily.

CatBoost is even more efficient version of XGBoost. It is open sourced by Yandex;

A model hyperparameter is a configuration that is external to the model and whose value cannot be estimated from data.
They are often used in processes to help estimate model parameters.
They are often specified by the practitioner.
They can often be set using heuristics.
They are often tuned for a given predictive modeling problem.

